

Streams & I/O

Luca Pozzoni, Giovanni Quattrocchi,
Danilo Ardagna

Politecnico di Milano
{nome.cognome}@polimi.it



Danilo Ardagna - Streams & I/O 2

Content

- Fundamental I/O concepts
- Files
 - Opening
 - Reading and writing streams
- File modes and Binary I/O
- String streams
 - Line-oriented input
- Examples

Danilo Ardagna - Streams & I/O 3

Files

- We turn our computers on and off
 - The contents of its main memory is transient
- Data needs to be preserved
 - It must be stored on disks and similar permanent storage devices
- A file is a sequence of bytes stored in permanent storage
 - A file has a name (identifier)
 - The data on a file has a format
- We can read/write a file if we know its name and format

Danilo Ardagna - Streams & I/O 4

A file

0 1 2




- At the fundamental level, a file is a sequence of bytes numbered from 0 upwards
- Other notions can be supplied by programs that interpret a "file format"
 - For example, the 6 bytes corresponding to "123.45" might be interpreted as the floating-point number 123.45

Danilo Ardagna - Streams & I/O 5

Files

- General model



Files

- To read a file
 - We must know its name
 - We must open it (for reading)
 - Then we can read it
 - Once finished, we must close it
 - That is typically done implicitly (when the stream object is destroyed)
- To write a file
 - We must name it
 - We must open it (for writing)
 - Or create a new file of that name
 - Then we can write it
 - Once finished, we must close it
 - That is typically done implicitly (when the stream object is destroyed)

Opening a file for reading

```
// ...
int main()
{
    cout << "Please enter input file name: ";
    string iname;
    cin >> iname;

    ifstream ist {iname}; // ifstream is an "input stream from a file"
                          // defining an ifstream with a name string
                          // opens the file of that name for reading

    if (!ist) error("can't open input file ", iname);
    // ...
}
```

input file stream

we have to check if there are errors in the opening

Opening a file for writing

```
// ...
cout << "Please enter name of output file: ";
string oname;
cin >> oname;

ofstream ofs {oname}; // ofstream is an "output stream from a file"
                      // defining an ofstream with a name string
                      // opens the file with that name for writing

if (!ofs) error("can't open output file ", oname);
// ...
}
```

notice that: if we mis-spell the name, another file will be created instead (≠ from the "ist")

Implicit close

- When an **fstream** object goes out of scope, the file it is bound to is automatically closed

```
if (read) {
    // create input and open the file
    ifstream input{name};
    if (input) { // if the file is ok, "process" this file
        process(input);
    } else {
        cerr << "couldn't open: " + name;
    } // input goes out of scope and is destroyed on each iteration (after the end of if(read))
}
```

Reading from a file

- Suppose a file contains a sequence of pairs representing hours and temperature readings
 - 0 60.7
 - 1 60.6
 - 2 60.3
 - 3 59.22
- The hours are numbered from 0 to 23
- No further format is assumed
 - Maybe we can do better than that (but not just now)
- Termination
 - Reaching the **end-of-file** terminates the read
 - Anything unexpected in the file terminates the read
 - E.g., character 'q'

Reading a file

```
struct Reading { // a temperature reading
    int hour; // hour after midnight [0,23]
    double temperature;
};

vector<Reading> temps; // create a vector to store the readings

int hour;
double temperature;
ifstream ist(fname);
while (ist >> hour >> temperature) { // read
    if (hour < 0 || 23 < hour)
        cout << "hour out of range" << endl;
    temps.push_back(Reading(hour, temperature)); // store
}
```

here we're actually missing the check if we can read the file (read/open)

this takes the first element of ist and puts it into "hour" then it puts the second element into "temperature" (space and end-of-row are automatically read)

No Copy or Assign for I/O Objects

- We cannot copy or assign objects of the IO types:

```
ofstream out1, out2;
out1 = out2; // error: cannot assign stream objects
ofstream print(ofstream); // error: can't initialize the ofstream param
out2 = print(out2); // error: cannot copy stream objects
```

Because we can't copy the IO types, we cannot have a parameter or return type that is one of the stream types

- Functions that do IO typically pass and return the stream through references
- Reading or writing an IO object changes its state, so the reference must not be const

```
void print(ofstream&); // OK
```

Use

```
void do_some_printing(Date d1, Date d2)
{
    cout << d1; // means:
    // operator<<(cout,d1);
    cout << d1 << d2;
    // means:
    // (cout << d1) << d2; same as:
    // (operator<<(cout,d1)) << d2; same as:
    // operator<<((operator<<(cout,d1)), d2);
}
```

call to the operator "<<" (insert) passing a reference to the standard output and a "Date"

File Modes and Binary I/O

File open modes

- By default, an ifstream opens its file for reading
- By default, an ofstream opens its file for writing.
- Alternatives:
 - `ios_base::app` // append (i.e., output adds to the end of the file)
 - `ios_base::ate` // "at end" (open and seek to end)
 - `ios_base::binary` // binary mode – beware of system specific behavior
 - `ios_base::in` // for reading
 - `ios_base::out` // for writing
 - `ios_base::trunc` // truncate file to 0-length
- A file mode is optionally specified after the name of the file:
 - `ofstream of1 {name1};` // defaults to `ios_base::out`
 - `ifstream if1 {name2};` // defaults to `ios_base::in`
 - `ofstream ofs {name, ios_base::app};` // append rather than overwrite
 - `fstream fs {"myfile", ios_base::in | ios_base::out};` // both in and out

One of the differences is that text files are variable in size depending on the size of the elements that we have

"123" → 3 bytes

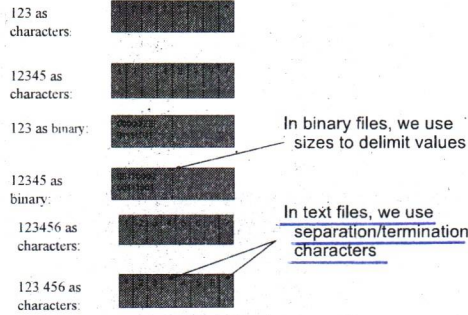
"12345" → 5 bytes

Meanwhile in binary:

"123" and "12345" occupies both 4 bytes.

However we want to avoid the binary open even if it occupies less memory.

Text vs. binary files



Text vs. binary

- Use text when you can
 - You can read it (without a fancy program)
 - You can debug your programs more easily
 - Text is portable across different systems
 - Most information can be represented reasonably as text
- Use binary when you must
 - E.g. image files, sound files

String Streams

String streams

- A **stringstream** reads/writes from/to a **string** rather than a file or a keyboard/screen

```
double str_to_double(string s)
// if possible, convert characters in s to floating-point value
{
    stringstream is(s); // make a stream so that
                        // we can read from s
    double d;
    is >> d;
    if (!is) error("double format error: ", s);
    return d;
}

double d1 = str_to_double("12.4"); // testing
double d2 = str_to_double("1.34e-3");
double d3 = str_to_double("twelve point three"); // will
                                                    // call error() !
```

input stringstream that we call "is": we initialize it with s

we define a double and we use the extract operator onto the stringstream (we perform the casting (conversion))

String streams

- See textbook for **ostringstream**
- String streams are very useful for
 - formatting into a fixed-sized space (think GUI)
 - for extracting typed objects out of a string

Type vs. line

Read a string

```
string name;
cin >> name; // input: Dennis Ritchie
cout << name << '\n'; // output: Dennis
```

Read a line

```
string name;
getline(cin, name); // input: Dennis Ritchie
cout << name << '\n'; // output: Dennis Ritchie
// now what?
// maybe:
istringstream ss(name);
ss >> first_name;
ss >> second_name;
```

Examples

Example 1: reading a CSV file

- An istringstream is often used when we have some work to do on an entire line, and other work to do with individual words within a line

```
Morgan,2015552368,8625550123
Drew,9735550130
Lee,6095550132,2015550175,8005550000
```

```
// members are public by default
struct PersonInfo {
string name;
vector<string> phones;
};
```

(why not int/double?)
When we read from text we may not want to cast the values if it's not necessary (moreover, we do it for robustness: maybe there's a "+33")

one element per line

```
vector<PersonInfo> people; // will hold all the records from the input
string line;
ifstream data("data.txt");
// read the input a line at a time until cin hits end-of-file (or another error)
while (getline(data, line)) {
    PersonInfo info; // create an object to hold this record's data
    istringstream record(line); // bind record to the line we just read
    // read the name
    if note that we are changing the delimiter of getline to ","
    getline(record, info.name, ',');
    string phone;
    // read the phone numbers
    while (getline(record, phone, ',')) {
        info.phones.push_back(phone); // and store them
    }
    people.push_back(info); // append this record to people
}
```

we call istringstream:
we create an object called "record" and we initialize it with "line"

notice that we have to do this instead of using directly "line" because we want a stringstream, while "line" is just a string

we read the whole "record" and we take everything before the first "," and we put it in "info.name"

("," = delimiter)

Example 2: A Word Transformation Map

Write a program that given one string, transforms it into another. The input to our program is two files. The first file contains **rules** that we will use to **transform** the **text** in the second file. Each rule consists of a word that might be in the input file and a phrase to use in its place.

```
word-transformation file:
y why
r are
u you
second file:
where r u
output file:
where are you
```

An example: A Word Transformation Map

```
void word_transform(istream &map_file, istream &input)
{
    auto trans_map = buildMap(map_file);
    string text;
    while (getline(input, text)) {
        istringstream stream(text);
        string word;
        bool firstword = true;
        while (stream >> word) {
            if (firstword)
                firstword = false;
            else
                cout << " ";
            cout << transform(word, trans_map);
        }
        cout << endl;
    }
}
```

An example: A Word Transformation Map

```
map<string, string> buildMap(istream &map_file)
{
    map<string, string> trans_map;
    string key;
    string value;
    while (map_file >> key && getline(map_file,
                                     value))
        if (value.size() > 1)
            trans_map[key] = value.substr(1);
        else
            cout << "no rule for " + key << endl;
    return trans_map;
}
```

An example: A Word Transformation Map

```
const string &
transform(const string &s, const map<string,
                                     string> &m)
{
    auto map_it = m.find(s);
    if (map_it != m.cend())
        // if this word is in the transformation map
        return map_it->second;
    else
        return s;
}
```

Readings

User-defined output: operator<<()

- Usually trivial:

```
ostream& operator<<( ostream& os, const Date& d )
{
    return os << '(' << d.year()
               << ',' << d.month()
               << ',' << d.day() << ')';
}
```

- We often use several different ways of outputting a value
 - Tastes for output layout and detail vary

User-defined input: operator>>()

```
istream& operator>>(istream& is, Date& dd)
// Read date in format: year month day
{
    int y, d, m;
    if (is >> y >> m >> d) {
        dd = Date{y,m,d}; // update dd
    }
    return is;
}
```

Binary files

```
int main() // use binary input and output
{
    cout << "Please enter input file name\n";
    string iname;
    cin >> iname;
    ifstream ifs (iname, ios_base::binary); // note: binary
    if (!ifs) error("can't open input file ", iname);

    cout << "Please enter output file name\n";
    string oname;
    cin >> oname;
    ofstream ofs (oname, ios_base::binary); // note: binary
    if (!ofs) error("can't open output file ", oname);

    // "binary" tells the stream not to try anything clever operation
    // with the bytes
}
```

Binary files

```
vector<int> v;

// read from binary file:
for (int i; ifs.read(as_bytes(i), sizeof(int)); )
// note: reading bytes
    v.push_back(i);

// ... do something with v ...

// write to binary file:
for (int i=0; i<v.size(); ++i)
    ofs.write(as_bytes(v[i]), sizeof(int)); // note: writing
// bytes


return 0;
}
```

// For now, treat as_bytes() as a primitive
// Warning! Beware transferring between different systems

Positioning in a filestream

- To support random access, the system maintains a marker that determines where the next read or write will happen
- We also have two functions:
 - One **repositions the marker** by seeking to a given position
 - The second **tells the current position** of the marker
- The library actually defines two pairs of seek and tell functions:
 - One pair is used by input streams, the other by output streams
 - The input and output versions are distinguished by a suffix that is either a **g** ("getting", i.e. reading data), or **p** ("putting", i.e. writing data)

Positioning in a filestream



```
fstream fs (name); // open for input and output
// ...
fs.seekg(5); // move reading position ('g' for 'get') to 5 (the 6th character)
char ch;
fs>>ch; // read the x and increment the reading position to 6
cout << "sixth character is " << ch << " (" << int(ch)
    << ")\n";
fs.seekp(1); // move writing position ('p' for 'put') to 1 (the 2nd character)
fs<<"y"; // write and increment writing position to 2
```

Positioning in a filestream

- We can use only the **g** versions on an **istream** and on the types that inherit from it, **ifstream** and **istreamstring**
- We can use only the **p** versions on an **ostream** and on the types that inherit from it, **ofstream** and **ostreamstring**
- An **istream**, **fstream**, or **stringstream** can both **read** and **write** the associated stream; we can use either the **g** or **p** versions on objects of these types

There Is Only One Marker

- The fact that the library distinguishes between the "putting" and "getting" versions of the **seek** and **tell** functions can be misleading
- Even though the library makes this distinction, it **maintains only a single marker in a stream** — there are no distinct **read** and **write** markers

Repositioning the Marker

```
// set the marker to a fixed position
seekg(new_position); // set the read marker to the given pos_type
// location
seekp(new_position); // set the write marker to the given pos_type
// location
// offset some distance ahead of or behind the given starting point
seekg(offset, from); // set the read marker offset distance from
// from seekp(offset, from); // offset has type off_type
```

Reading and writing to the same file

```
abcd
efg
hi
j
      →
abcd
efg
hi
j
5 9 12 14
```

Reading and writing to the same file

```
int main()
{
    // open for input and output and preposition file pointers to end-of-file
    // file mode argument
    fstream inOut("copyOut", fstream::ate | fstream::in |
                  fstream::out);
    if (!inOut) {
        cerr << "Unable to open file!" << endl;
        return EXIT_FAILURE; // EXIT_FAILURE
    }
    // inOut is opened in ate mode, so it starts out positioned at the end
    auto end_mark = inOut.tell(); // remember original end-of-file
    // position
    inOut.seekg(0, fstream::beg); // reposition to the start of the file
    size_t cnt = 0; // accumulator for the byte count string line;
    string line; // hold each line of input
```


Reading and writing to the same file

```
// while we haven't hit an error and are still reading the original
// data and can get another line of input
while (inOut && inOut.tellg() != end_mark
      && getline(inOut, line))
{
    cnt += line.size() + 1; // add 1 to account for the newline
    auto mark = inOut.tellg(); // remember the read position
    inOut.seekp(0, fstream::end); // set the write marker to the end
    inOut << cnt; // write the accumulated length
    // print a separator if this is not the last line
    if (mark != end_mark) inOut << " ";
    inOut.seekg(mark); // restore the read position
}
inOut.seekp(0, fstream::end); // seek to the end
inOut << "\n"; // write a newline at end-of-file
return 0;
}
```

Positioning

- Whenever you can
 - Use simple streaming
 - Streams/streaming is a very powerful metaphor
 - Write most of your code in terms of "plain" istream and ostream
 - Positioning is far more error-prone
 - Handling of the end of file position is system dependent and basically unchecked

Using ostringstreams

- An **ostringstream** is useful when we need to build up our output a little at a time but **do not** want to print the output until later
- For example, we might want to validate and reformat the phone numbers we read in the previous example
 - If all the numbers are valid, we want to print a new file containing the reformatted numbers
 - If a person has any invalid numbers, we won't put them in the new file. Instead, we'll write an error message containing the person's name and a list of their invalid numbers

Using ostringstreams

```
for (const auto &entry : people) { // for each entry in people
    ostringstream formatted, badNums; // objects created on each loop
    for (const auto &nums : entry.phones) { // for each number
        if (!valid(nums)) {
            badNums << " " << nums; // string in badNums
        } else
            // "writes" to formatted's string
            formatted << " " << format(nums);
    }
    if (badNums.str().empty()) // there were no bad numbers
        os << entry.name << " " // print the name
        << formatted.str() << endl; // and reformatted numbers
    else
        // otherwise, print the name and bad numbers
        cerr << "input error: " << entry.name
        << " invalid number(s) " << badNums.str() << endl;
}
```

References

- Lippman Chapters 8, 17

Credits

• Bjarne Stroustrup. www.stroustrup.com/Programming