

```

//-----
// MPI - MESSAGE PASSING INTERFACE
//-----

#include<iostream>
#include<mpi.h>

int main(int argc, char *argv[]){
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);    // size = #processors
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);    // rank = ID of the processor
    .. ;
    MPI_Finalize();
    return 0;
}

// * if we pass 0 arguments -> argc=1, argv[0] is the name of the program
// * if we pass k arguments -> argc=k+1, argv[1] first parameter, .., argv[k] k-th input

// COMPILE and RUN
$ mpicxx --std=c++11 file_name.cc -o exe_name           // 1.
$ mpicxx --std=c++11 file_name.cc                     // 2.
$ mpicxx --std=c++11 file_name.cc file1.cc file2.cc -o exe_name // 3. (e.g. with classes)
$ mpicxx --std=c++11 file_name.cc file1.cc file2.cc     // 4.

$ mpiexec -np=4 exe_name                               // 1.a, 3.a
$ mpiexec -np=4 exe_name text1 text2                   // 1.b, 3.b
$ mpiexec -np=4 a.out                                  // 2.a, 4.a
$ mpiexec -np=4 a.out text1 text2                      // 2.b, 4.b
// * 1.a, 2.a, 3.a, 4.a: running on 4 processors, argc=1, argv[0] is the name of the program
// * 1.b, 2.b, 3.b, 4.b: running on 4 processors, argc=3, argv[1]="text1", argv[2]="text3"

//-----
// COMMUNICATORS
//-----
// * point-to-point: involves two processors -> sender, receiver
// * collective : involves all processors -> broadcast, reduce, scatter, gather,

// *****
// SEND and RECEIVE
// *****

MPI_Send(const void *buf,                // where is what we want to send?
         int count, MPI_Datatype datatype, // how many (at most) and of what type?
         int dest, int tag, MPI_Comm comm) // to who? tag? through what?

MPI_Recv(void *buf,                      // where do we store?
         int count, MPI_Datatype datatype, // how many (at most) and of what type?
         int source, int tag, MPI_Comm comm, // from who? tag? through what?
         MPI_Status *status)              // "MPI_STATUS_IGNORE"

// Example: rank0 takes as input a double and sends it to the others
double n;
if(rank==0){
    cin >> n;
    for(int dest=1; dest<size; ++dest)
        MPI_Send(&n, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
}
else{
    MPI_Recv(&n, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}

// *****
// BROADCAST
// *****
// Rank_0 sends a message to all the other ranks, without doing it one by one:
// * cannot use tags
// * every other process in the communicator MUST call the collective function

MPI_Bcast(void *buffer,                  // sender: where is stored the info
          int count, MPI_Datatype datatype, // receiver: where to store the info
          int root, MPI_Comm comm)       // how many (at most) and of what type?
                                             // who owns the data? source. Through what?

// Example: rank0 takes as input a double and sends it to the others
double n;
if(rank==0)
    cin >> n;
MPI_Bcast(&n, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);

// *****
// REDUCE/ALL_REDUCE
// *****
// Perform efficiently some operations (max, min, sum, prod) among ranks.

MPI_Reduce(const void *sendbuf, void *recvbuf, // send buffer, receiver buffer
           int count, MPI_Datatype datatype, // how many (at most) and of what type?
           MPI_Op op, int dest, MPI_Comm comm) // what operation? send to who? through what?

// Example: every rank has a local value and we want the sum of them in rank_0
// Note: we declare "total" in each rank, not only in the destination one.
double local_value;
double total;
MPI_Reduce(&local_value, &total, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

```

```

// What if we want the sum to be saved in the variable total of each rank?

MPI_Allreduce(const void *sendbuf, void *recvbuf, // send buffer, receiver buffer
              int count, MPI_Datatype datatype, // how many (at most) and of what type?
              MPI_Op op, MPI_Comm comm) // what operation? through what?

// MPI_IN_PLACE
// With MPI_IN_PLACE we can use a single buffer for both input and output.
// (we cannot do, for example: MPI_Allreduce(&minimum, &minimum, ... ); )
double minimum;
MPI_Allreduce(MPI_IN_PLACE, &minimum, 1, MPI_DOUBLE, MPI_MIN, MPI_COMM_WORLD);

// *****
// SCATTER and GATHER
// *****
// When we have to divide data between ranks we can proceed in two ways:
// CYCLIC PARTITIONING: used when data source is available across all processes

for(size_t i=rank; i<v.size(); i+=size) { .. }

// BLOCK PARTITIONING: used when data source is available on a single process (#elems % #process = 0)

MPI_Scatter(const void *sendbuf, // send buffer
            int sendcount, MPI_Datatype sendtype, // #elems we send to individual rank? type?
            void *recvbuf, // receiver buffer
            int recvcnt, MPI_Datatype recvttype, // #elems will be received? type?
            int root, MPI_Comm comm) // source of the data? through what?

// The dual of MPI_Scatter (decomposes one into many) is MPI_Gather (composes many into one):

MPI_Gather(const void *sendbuf, // send buffer
            int sendcount, MPI_Datatype sendtype, // #elems we send? type?
            void *recvbuf, // receiver buffer
            int recvcnt, MPI_Datatype recvttype, // #elems will be received? type?
            int root, MPI_Comm comm) // destination of the data? through what?

// If we want the send buffer to be update in all ranks, then:

MPI_Allgather(const void *sendbuf, // send buffer
              int sendcount, MPI_Datatype sendtype, // #elems we send? type?
              void *recvbuf, // receiver buffer
              int recvcnt, MPI_Datatype recvttype, // #elems will be received? type?
              MPI_Comm comm) // through what?

//-----
// READ_VECTOR and PRINT_VECTOR
//-----

vector<double> x = mpi::read_vector(n, "x", MPI_COMM_WORLD);
vector<double> y = mpi::read_vector(n, "y", MPI_COMM_WORLD);
vector<double> z = x+y;
mpi::print_vector(z, n, "The result is ..", MPI_COMM_WORLD);

// mpi::read_vector : read and spread the informations among processes (MPI_Scatter)
// mpi::print_vector: build a single vector from Locals and print it (MPI_Gather, rank_0)

//-----
// Random REMEMBER
//-----
// 1. When inside the code we have that rank_0 takes something in input (e.g. cin>>a>>b>>c;) we do:
$ mpicxx --std=c++11 file_name.cc
$ mpiexec -np 4 a.out
5 6 7

// In this way we perform the assignment: a=5, b=6, c=7. More fast:
$ vim input_short // creates a file where we write "5 6 7"
$ cat input_short // opens the file
5 6 7
$ mpiexec -np 4 a.out < input_short // the "<" means "take inputs from .."
$ mpiexec -np 4 a.out < input_short > output_short // output saved in "output_short" (nothing displayed)

// 2. * STANDARD INPUT: if we have std::cin (rank_0) in the code -> separate line of inputs in the cmd line
// (after this rank_0 has to pass the informations through broadcast/scatter/send)
$ mpicxx --std=c++11 file_name.cc
$ mpiexec -np 4 a.out
5 6 7 // cin>>a>>b>>c

// * USER PASSES INFORMATIONS: if we get something from the arguments -> we write arguments in the cmd line
// (after this the information is available for all the processes)
$ mpicxx --std=c++11 file_name.cc
$ mpiexec -np 4 a.out 5 6 7

// ALL processors have the informations, but we need conversions (argv is a vector of strings).
if(argc == 4){ // we have arguments -> argc = 1 + #arguments
    unsigned var1 = std::stoul(argv[1]); // var1 = 5; (unsigned)
    unsigned var2 = std::stoul(argv[2]); // var2 = 6; (unsigned)
    int var3 = std::stoi(argv[3]); // var3 = 7; (integer)
}

// 3. We can use MPI_Reduce/MPI_Allreduce also on vectors
int N = 10;
vector<double> local_x(N), sum(N);
MPI_Reduce(local_x.data(), sum.data(), N, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
for vectors and matrices
MPI_Bcast(&the_string[0], length, MPI_CHAR, 0, MPI_COMM_WORLD);

```



```

#include<iostream>
#include<mpi.h>

int main(int argc, char* argv[]){

    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    .. ;

    MPI_Finalize();
    return 0;
}

// Send/Receive (point-to-point)
MPI_Send(&n, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
MPI_Recv(&n, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);

// Broadcast (collective)
MPI_Bcast(&n, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);

// Reduce/Allreduce (MPI_MAX, MPI_MIN, MPI_SUM, MPI_PROD)
MPI_Reduce(&local_value, &total, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
MPI_Allreduce(MPI_IN_PLACE, &value, 1, MPI_DOUBLE, MPI_MIN, MPI_COMM_WORLD);

// Cyclic partitioning
for(size_t i=rank; i<v.size(); i+=size){ .. }

// Block partitioning: Scatter/Gather
MPI_Scatter(global.data(), local_n, MPI_DOUBLE, local.data(), local_n, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Gather(local.data(), local_n, MPI_DOUBLE, global.data(), local_n, MPI_DOUBLE, 0, MPI_COMM_WORLD);
MPI_Allgather(local.data(), local_n, MPI_DOUBLE, global.data(), local_n, MPI_DOUBLE, MPI_COMM_WORLD);

// Read_vector/Print_vector
vector<double> x = mpi::read_vector(n, "x", MPI_COMM_WORLD);
vector<double> y = mpi::read_vector(n, "y", MPI_COMM_WORLD);
vector<double> z = x+y;
mpi::print_vector(z, n, "The global vector is: ", MPI_COMM_WORLD);

//-----

$ mpicxx --std=c++11 file_name.cc
$ mpiexec -np=4 a.out

// Name specification
$ mpicxx --std=c++11 file_name.cc -o exe_name
$ mpiexec -np=4 exe_name

// More files
$ mpicxx --std=c++11 file_name.cc file1.cc file2.cc -o exe_name
$ mpiexec -np=4 exe_name

// Arguments
$ mpicxx --std=c++11 file_name.cc -o exe_name
$ mpiexec -np=4 exe_name text1 text2          // argc=3, argv[1]="text1", argv[2]="text2"

// Standard input (only rank_0 has "std::cin >> a >> b >> c;")
$ mpicxx --std=c++11 file_name.cc -o exe_name
$ mpiexec -np=4 exe_name
5 6 7          // a=5, b=6, c=7

// Standard input with saved input and output
$ vim input_short
$ mpicxx --std=c++11 file_name.cc -o exe_name
$ mpiexec -np=4 exe_name < input_short > output_short

// User passed informations
$ mpicxx --std=c++11 file_name.cc -o exe_name
$ mpiexec -np=4 exe_name 5 6 7
// ..
if (argc==4){
    unsigned var1 = std::stoul(argv[1]);
    unsigned var2 = std::stoul(argv[2]);
    int var3 = std::stoi(argv[3]);
}

```