

Quickstart to torchdyn

`torchdyn` is a PyTorch library dedicated to neural differential equations and equilibrium models.

Central to the `torchdyn` approach are continuous and implicit neural networks, where *depth* is taken to its infinite limit.

This notebook serves as a gentle introduction to NeuralODE, concluding with a small overview of `torchdyn` features.

```
[1]: from torchdyn.core import NeuralODE
      from torchdyn.datasets import *
      from torchdyn import *

      %load_ext autoreload
      %autoreload 2

[2]: # quick run for automated notebook validation
      dry_run = False
```


Generate data from a static toy dataset

We'll be generating data from toy datasets. In `torchdyn`, we provide a wide range of datasets often used to benchmark and understand Neural ODEs. Here we will use the classic moons dataset and train a Neural ODE for binary classification

```
[3]: d = ToyDataset()
      X, yn = d.generate(n_samples=512, noise=1e-1, dataset_type='moons')

[4]: import matplotlib.pyplot as plt

      colors = ['orange', 'blue']
      fig = plt.figure(figsize=(3,3))
      ax = fig.add_subplot(111)
      for i in range(len(X)):
          ax.scatter(X[i,0], X[i,1], s=1, color=colors[yn[i].int()])
```



Generated data can be easily loaded in the dataloader with standard `PyTorch` calls

```
[5]: import torch
      import torch.utils.data as data
      device = torch.device("cpu") # all of this works in GPU as well :)

      X_train = torch.Tensor(X).to(device)
      y_train = torch.LongTensor(yn.long()).to(device)
      train = data.TensorDataset(X_train, y_train)
      trainloader = data.DataLoader(train, batch_size=len(X), shuffle=True)
```

We utilize Pytorch Lightning to handle training loops, logging and general bookkeeping. This allows `torchdyn` and Neural Differential Equations to have access to modern best practices for training and experiment reproducibility.

In particular, we combine modular `torchdyn` models with `LightningModules` via a `Learner` class:

```
[6]: import torch.nn as nn
      import pytorch_lightning as pl

      class Learner(pl.LightningModule):
          def __init__(self, t_span:torch.Tensor, model:nn.Module):
              super().__init__()
              self.model, self.t_span = model, t_span

          def forward(self, x):
```

```

    return self.model(x)

def training_step(self, batch, batch_idx):
    x, y = batch
    t_eval, y_hat = self.model(x, t_span)
    y_hat = y_hat[-1] # select last point of solution trajectory
    loss = nn.CrossEntropyLoss()(y_hat, y)
    return {'loss': loss}

def configure_optimizers(self):
    return torch.optim.Adam(self.model.parameters(), lr=0.01)

def train_dataloader(self):
    return trainloader

```

Define a Neural ODE

Analogously to most forward neural models we want to realize a map

$$x \mapsto \hat{y}$$

where \hat{y} becomes the best approximation of a true output y given an input x . In torchdyn you can define very simple Neural ODE models of the form

$$\begin{cases} \dot{z}(t) = f(z(t), \theta) \\ z(0) = x \\ \hat{y} = z(1) \end{cases} \quad t \in [0, 1]$$

by just specifying a neural network f and giving some simple settings.

Note: This Neural ODE model is of *depth-invariant* type as neither f explicitly depend on s nor the parameters θ are depth-varying. Together with their *depth-variant* counterpart with s concatenated in the vector field was first proposed and implemented by [Chen T. Q. et al, 2018]

Define the vector field (DEFunc)

The first step is to define any PyTorch `torch.nn.Module`. This takes the role of the Neural ODE vector field $f(h, \theta)$

```

[16]: f = nn.Sequential(
        nn.Linear(2, 16),
        nn.Tanh(),
        nn.Linear(16, 2)
    )
    t_span = torch.linspace(0, 1, 5)

```

In this case we chose f to be a simple MLP with one hidden layer and `tanh` activation

Define the NeuralDE

The final step to define a Neural ODE is to instantiate the torchdyn's class `NeuralDE` passing some customization arguments and `f` itself.

In this case we specify: * we compute backward gradients with the `'adjoint'` method. * we will use the `'dopri5'` (Dormand-Prince) ODE solver from `torchdyn`, with no additional options;

```

[17]: model = NeuralODE(f, sensitivity='adjoint', solver='dopri5').to(device)

Your vector field callable (nn.Module) should have both time 't' and state 'x' as
arguments, we've wrapped it for you.

```

Train the Model

With the same forward method of `NeuralDE` objects you can quickly evaluate the entire trajectory of each data point in `x_train` on an interval `t_span`

```

[18]: t_span = torch.linspace(0, 1, 100)
    t_eval, trajectory = model(x_train, t_span)
    trajectory = trajectory.detach().cpu()

```

The numerical method used to solve a `NeuralODE` have great effect on its speed. Try retraining with the following

```

[23]: f = nn.Sequential(
        nn.Linear(2, 16),

```

```

        nn.Tanh(),
        nn.Linear(16, 2)
    )

model = NeuralODE(f, sensitivity='adjoint', solver='rk4', solver_adjoint='dopri5',
    atol_adjoint=1e-4, rtol_adjoint=1e-4).to(device)

learn = Learner(t_span, model)
if dry_run: trainer = pl.Trainer(min_epochs=1, max_epochs=1)
else: trainer = pl.Trainer(min_epochs=200, max_epochs=300)
trainer.fit(learn)

GPU available: True, used: False
TPU available: False, using: 0 TPU cores

| Name | Type | Params
-----
0 | model | NeuralODE | 82
-----
82 | Trainable params
0 | Non-trainable params
82 | Total params
0.000 | Total estimated model params size (MB)

Your vector field callable (nn.Module) should have both time 't' and state 'x' as
arguments, we've wrapped it for you.

```

Plot the Training Results

We can first plot the trajectories of the data points in the depth domain s

```


[24]: t_eval, trajectory = model(X_train, t_span)
trajectory = trajectory.detach().cpu()

[25]: color=['orange', 'blue']

fig = plt.figure(figsize=(10,2))
ax0 = fig.add_subplot(121)
ax1 = fig.add_subplot(122)
for i in range(500):
    ax0.plot(t_span, trajectory[:,i,0], color=color[int(yn[i])], alpha=.1);
    ax1.plot(t_span, trajectory[:,i,1], color=color[int(yn[i])], alpha=.1);
ax0.set_xlabel(r"$t$ [Depth]"); ax0.set_ylabel(r"$h_0(t)$")
ax1.set_xlabel(r"$t$ [Depth]"); ax1.set_ylabel(r"$z_1(t)$")
ax0.set_title("Dimension 0"); ax1.set_title("Dimension 1")

[25]: Text(0.5, 1.0, 'Dimension 1')

```



Then the trajectory in the *state-space*


As you can see, the Neural ODE steers the data-points into regions of null loss with a continuous flow in the depth domain. Finally, we can also plot the learned vector field f

```

[26]: # evaluate vector field
n_pts = 50
x = torch.linspace(trajectory[:,0].min(), trajectory[:,0].max(), n_pts)
y = torch.linspace(trajectory[:,1].min(), trajectory[:,1].max(), n_pts)
X, Y = torch.meshgrid(x, y); z = torch.cat([X.reshape(-1,1), Y.reshape(-1,1)], 1)
f = model.vf(z.to(device)).cpu().detach()
fx, fy = f[:,0], f[:,1]; fx, fy = fx.reshape(n_pts, n_pts), fy.reshape(n_pts, n_pts)
# plot vector field and its intensity
fig = plt.figure(figsize=(4, 4)); ax = fig.add_subplot(111)
ax.streamplot(X.numpy().T, Y.numpy().T, fx.numpy().T, fy.numpy().T, color='black')
ax.contourf(X.T, Y.T, torch.sqrt(fx.T**2+fy.T**2), cmap='RdYlBu')

[26]: <matplotlib.contour.QuadContourSet at 0x7f11ed2d5e50>

```



Sweet! You trained your first Neural ODEs! Now you can proceed and learn about more advanced models with the next tutorials

More about `torchdyn`

```
[46]: import time
      from torchdyn.numerics import Euler, RungeKutta4, Tsitouras45, DormandPrince45, MSZero,
      Euler, HyperEuler
      from torchdyn.numerics import odeint, odeint_mshooting, Lorenz
      from torchdyn.core import ODEProblem, MultipleShootingProblem
```

But wait! `torchdyn` has way more than `NeuralODEs`. If you wish to solve generic differential equations parallelizable both in space (initial conditions) as well in time, with `parallel`, but do not need neural networks inside the vector field, you can use our functional API like so:

```
[47]: x0 = torch.randn(8, 3) + 15
      t_span = torch.linspace(0, 3, 2000)
      sys = Lorenz()

[48]: t0 = time.time()
      t_eval, accurate_sol = odeint(sys, x0, t_span, solver='dopri5', atol=1e-6, rtol=1e-6)
      accurate_sol_time = time.time() - t0

      t0 = time.time()
      t_eval, base_sol = odeint(sys, x0, t_span, solver='euler')
      base_sol_time = time.time() - t0

      t0 = time.time()
      t_eval, rk4_sol = odeint(sys, x0, t_span, solver='rk4')
      rk4_sol_time = time.time() - t0

      t0 = time.time()
      t_eval, dp5_low_sol = odeint(sys, x0, t_span, solver='dopri5', atol=1e-3, rtol=1e-3)
      dp5_low_time = time.time() - t0

      t0 = time.time()
      t_eval, ms_sol = odeint_mshooting(sys, x0, t_span, solver='mszero', fine_steps=2,
      maxiter=4)
      ms_sol_time = time.time() - t0
```

Alternatively, you can wrap your vector field in a specific `*Problem` to perform sensitivity analysis and optimize for terminal as well as integral objectives:

```
[52]: prob = ODEProblem(sys, sensitivity='interpolated_adjoint', solver='dopri5', atol=1e-3,
      rtol=1e-3,
      solver_adjoint='tsit5', atol_adjoint=1e-3, rtol_adjoint=1e-3)
      t0 = time.time()
      t_eval, sol_torchdyn = prob.odeint(x0, t_span)
      t_end1 = time.time() - t0
```

Our numerics suite includes other tools, such as a `odeint_hybrid` for hybrid systems (potentially stochastic and multi-mode). We have built our numerics suite from the ground up to be compatible with hybridized methods such as hypersolvers, where a base solver works in tandem with neural approximators to increase accuracy while retaining improved extrapolation capabilities. In fact, these methods can be called from the same API:

```
[40]: class VanillaHyperNet(nn.Module):
      def __init__(self, net):
          super().__init__()
          self.net = net
          for p in self.net.parameters():
              torch.nn.init.uniform_(p, 0, 1e-5)
      def forward(self, t, x):
          return self.net(x)

[43]: net = nn.Sequential(nn.Linear(3, 64), nn.Softplus(), nn.Linear(64, 64), nn.Softplus(),
      nn.Linear(64, 3))
      hypersolver = HyperEuler(VanillaHyperNet(net))
      t_eval, sol = odeint(sys, x0, t_span, solver=hypersolver) # note: this has to be trained!
```

We also provide an extensive set of tutorial subdivided into modules. Each tutorial deals with a specific aspect of continuous or implicit models, or showcases applications (control, generative modeling, forecasting, optimal control of ODEs and PDEs, graph node classification). Check `torchdyn/tutorials` for more information.