

- **Bayesian linear regression and Ridge regression**

In the Bayesian linear regression settings, when $\mathbf{w}_0 = \mathbf{0}$ and $S_0 = \tau^2 \mathbb{I}$, we have:

$$\ln p(\mathbf{w}|\mathbf{t}) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - \mathbf{w}^t \phi(\mathbf{x}_n))^2 - \frac{1}{2\tau^2} \|\mathbf{w}\|_2^2$$

In this case, MAP, maximum a-posteriori, (\mathbf{w}_N) is equivalent to the solution of ridge regression ($\hat{\mathbf{w}}_{ridge}$) with $\lambda = \frac{\sigma^2}{\tau^2}$. Bayesian regression also allows to account for regularization (which can be obtained as a specific case).

- **k-Nearest Neighbor**

To classify a point \mathbf{x} : select the k most similar points to \mathbf{x} in the training set and assign to \mathbf{x} the most frequent class among those. No actual model is computed: we use the training records to predict an unknown class label. The elements are: the training dataset, the similarity function and the value of k .

1. Compute distance to other training records
2. Identify the k nearest neighbors
3. Use labels of nearest neighbors to determine the label of the point (majority voting/others)

How many neighbors? It's an hyperparameter (cross-validation changing k at every iteration or nested cross-validation). If k is too small the classification might be too sensitive. If k is too large the neighborhood may include quite dissimilar points.

consistent learners: $N \geq \frac{1}{\epsilon} \left(8 VC(H) \log_2 \left(\frac{13}{\epsilon} \right) + 4 \log_2 \left(\frac{2}{\delta} \right) \right)$

general learners:
with prob $\geq (1-\delta)$
every $h \in H$ satisfies $error_{true}(h) \leq error_D(h) + \sqrt{\frac{VC(H) \left(\ln \left(\frac{2N}{VC(H)} + 1 \right) + \ln 4/\delta \right)}{N}}$

- MDP: agent-environment interface (well defined states, actions, rewards)
- MDP: Bellman equations (expectation, optimality) to evaluate and find the optimal value functions/ policy
- Computationally unfeasible
- Dynamic programming: approximate the optimal value function/ policy iteratively (policy iteration, value iteration)
- Real application: unknown states, actions, rewards
- Monte Carlo methods: learn value functions/ policy through experience (data: complete episodes of states, actions, rewards)
 - policy evaluation for the evaluation of V_π
 - policy improvement? V_π useless, need for Q_π
 - Q_π evaluation: need for exploration
 - exploring starts, ϵ -soft

	Bellman expectation equation	Bellman optimality equation
Dynamic programming	• policy iteration	• value iteration
Empirical version	• Monte Carlo control • SARSA	• Q-learning

actual class :

		1	0
predicted :	1	tp	fp
	0	fn	tn

- accuracy: $Acc = \frac{tp+tn}{N}$

- precision: $Pre = \frac{tp}{tp+fp}$

$$\frac{\square}{\square \square}$$

- recall: $Rec = \frac{tp}{tp+fn}$

$$\frac{\square}{\square}$$

- F_1 score: $F_1 = \frac{2 \cdot Prec \cdot Rec}{Prec + Rec}$

SVM

Separable

minimize $\frac{1}{2} \|\underline{w}\|_2^2$
s.t. $t_n (\underline{w}^T \phi(x_n) + b) \geq 1$

maximize $\sum_n \alpha_n - \frac{1}{2} \sum_n \sum_m \alpha_n \alpha_m t_n t_m k(x_n, x_m)$
s.t. $\alpha_n \geq 0$
 $\sum \alpha_n t_n = 0$

$$y(x) = \sum_n \alpha_n t_n k(x, x_n) + b$$

$$b = \frac{1}{|S|} \sum_{x_n \in S} (t_n - \sum_{x_m \in S} \alpha_m t_m k(x_n, x_m))$$

Non-separable

minimize $\frac{1}{2} \|\underline{w}\|_2^2 + C \sum_n \xi_n$ (penalties to the violation)
s.t. $t_n (\underline{w}^T \phi(x_n) + b) \geq 1 - \xi_n$
 $\xi_n \geq 0$

maximize $-\frac{1}{2} \sum_n \sum_m \alpha_n \alpha_m t_n t_m k(x_n, x_m)$
s.t. $0 \leq \alpha_n \leq \frac{1}{N}$
 $\sum \alpha_n t_n = 0$
 $\sum \alpha_n \geq \nu$

parameter that controls both the margin errors and the # support vectors:

$$0 \leq \% \text{ margin errors} \leq \nu \leq \% \text{ \# support vectors} \leq 1$$